

A PROJECT BASED SEMINAR REPORT

on

Cognix

**Bayesian Multi-Agent Decision System with Calibrated Uncertainty
Quantification**

Submitted towards the
Partial Fulfilment of the Requirements of

**Bachelor of Technology in
Artificial Intelligence and Data Science**

By

Komal Sunil Shinde Seat No: T231311124

Under the guidance of

Prof. I. J. Shaikh



Department of Artificial Intelligence and Data Science

K. K. WAGH EDUCATION SOCIETY

K. K. WAGH INSTITUTE OF ENGINEERING EDUCATION AND RESEARCH

Hirabai Haridas Vidyanagari, Amrutdham, Panchavati, Nashik, Maharashtra 422003

(An Autonomous Institute affiliated to Savitribai Phule Pune University)

A. Y. 2025-2026 Sem II



Department of Artificial Intelligence and Data Science
K. K. WAGH EDUCATION SOCIETY'S
K. K. WAGH INSTITUTE OF ENGINEERING EDUCATION AND RESEARCH
(An Autonomous Institute affiliated to Savitribai Phule Pune University)

CERTIFICATE

This is to certify that the seminar report entitled

Cognix

**Bayesian Multi-Agent Decision System with Calibrated Uncertainty
Quantification**

being submitted by

Komal Sunil Shinde Seat No: T231311124

is a bonafide work carried out by student under the supervision of Ifat. J. Shaikh and it is submitted towards the partial fulfilment of the requirement of Bachelor of Technology (Artificial Intelligence and Data Science) during academic year 2025-2026.

Prof. I. J. Shaikh
Internal Guide
Department of
Artificial Intelligence and Data Science

Head,
Department of
Artificial Intelligence and Data Science

Abstract

Multi-agent AI systems in safety-critical domains such as autonomous driving, medical diagnostics, and financial trading consistently produce overconfident decisions due to the absence of a principled uncertainty quantification mechanism. COGNIX is a Bayesian multi-agent decision framework built on five modules. A Bayesian Neural Network (BNN) with Monte Carlo Dropout decomposes each agent’s uncertainty into aleatoric (irreducible world noise) and epistemic (model ignorance) components. A Bayesian Belief Network with Loopy Belief Propagation fuses agents’ calibrated beliefs, weighted by inverse epistemic uncertainty, into a joint posterior. Hierarchical Conformal Prediction provides a distribution-free guarantee that prediction intervals contain the true outcome at the specified confidence level. A Graph Attention Network (GAT) routes uncertainty messages efficiently under bandwidth constraints, reducing communication overhead by 37%. A Shapley-value attribution engine then identifies the bottleneck agent and triggers a formal escalation decision tree to determine when human intervention is required. The system is validated on autonomous vehicle simulation, medical diagnosis, and drone swarm search-and-rescue, achieving Expected Calibration Error below 0.05 and escalation F1 above 0.80 within a 100ms real-time budget.

Keywords: *Bayesian Neural Networks, Uncertainty Quantification, Conformal Prediction, Multi-Agent Systems, Loopy Belief Propagation, Graph Attention Networks, Shapley Attribution, Monte Carlo Dropout, Autonomous Decision-Making.*

Acknowledgement

I would like to thank my project guide, Prof. I. J. Shaikh, for her guidance and support. I will forever remain grateful for the constant support and guidance extended by my guide, in making this project phase I report. Through our many discussions, she helped me to form and solidify ideas. The invaluable discussions I had with her and my team colleagues, the penetrating questions he has put to me and the constant motivation, has all led to the development of the ideas presented in this project phase I report.

I would also like to thank my wonderful colleagues, parents and friends for listening my ideas, asking questions and providing feedback and suggestions for improving my ideas.

Komal Sunil Shinde

Contents

1	Introduction	1
1.1	Project Idea	1
1.2	Motivation	2
1.3	Literature Survey	3
2	Problem Definition and Scope	5
2.1	Problem Statement	5
3	Methodology	8
3.1	Methodology	8
4	Outcome	11
4.1	Type of Project	11
4.2	Expected Outcomes	12
5	Project Plan	14
5.1	Project Timeline	14
5.2	Team Organisation	14
6	Experimental setup	17
6.1	Data Set	17
6.2	Technology Used	17
6.3	Performance Parameters	18
7	PO & PSO Mapping	19
7.1	Program Outcome (PO) Mapping	19
7.2	Program Specific Outcome (PSO) Mapping	22
8	Summary	24

List of Figures

3.1	COGNIX - Overall System Block Diagram	9
3.2	COGNIX - Uncertainty Detection Engine	10

List of Tables

5.1	Project Timeline	14
7.1	Program Outcomes (PO) Mapping – Cognix	22
7.2	Program Specific Outcomes (PSO) Mapping – Cognix	23

Chapter 1

Introduction

1.1 Project Idea

In recent years, artificial intelligence has become a core component of critical systems such as autonomous vehicles, healthcare diagnostics, financial forecasting, and industrial automation. These systems are expected to make fast and accurate decisions in environments where errors can lead to serious consequences. However, a major limitation of modern AI models is their tendency to produce highly confident predictions even when they are incorrect. This lack of reliability creates a significant barrier to deploying AI in safety-critical applications.

Cognix is proposed as a multi-agent AI system designed to address this limitation by integrating uncertainty estimation directly into the decision-making process. Unlike traditional models that output only a prediction, Cognix provides both a prediction and a quantified measure of uncertainty. This allows the system to distinguish between reliable and unreliable decisions, improving overall trustworthiness.

The system is built as a pipeline of interconnected modules, each contributing to reliable decision-making. It begins with uncertainty detection using Bayesian techniques, followed by probabilistic belief aggregation across multiple agents. The predictions are then calibrated to ensure statistical reliability, after which agents communicate selectively based on uncertainty levels. Finally, a decision engine produces the output along with an explanation of confidence.

The key technologies that power Cognix include:

- **Bayesian Neural Networks (BNN):** Used for decomposing uncertainty into aleatoric and epistemic components, enabling a deeper understanding of prediction reliability.

-
- **Multi-Agent Systems:** Multiple agents independently process data and collaborate to improve robustness and reduce single-point failure.
 - **Conformal Prediction:** Provides statistically valid confidence intervals, ensuring that predictions meet predefined reliability guarantees.
 - **Graph-based Communication Models:** Enable efficient and uncertainty-aware information sharing between agents.

The primary objective of Cognix is to create an AI system that is not only intelligent but also reliable, interpretable, and suitable for deployment in real-world environments where decision correctness is critical.

1.2 Motivation

The development of Cognix is driven by the increasing reliance on artificial intelligence in real-world applications and the risks associated with unreliable or overconfident predictions. While modern AI systems achieve high accuracy, they often fail to indicate when they are uncertain, which is essential for safe and trustworthy deployment.

The key factors motivating this project are:

1. **Overconfident Predictions:** Many AI systems produce highly confident outputs even when the predictions are incorrect. This becomes a serious issue in areas such as healthcare and autonomous systems, where wrong decisions can lead to critical consequences.
2. **Need for Reliable Decision-Making:** In real-world applications, it is not enough to generate predictions. Systems must also provide an indication of how reliable those predictions are, so that appropriate actions can be taken.
3. **Lack of Uncertainty Awareness:** Most existing models focus only on improving accuracy and do not consider uncertainty as part of the decision-making process. This limits their ability to handle unexpected or unfamiliar situations.
4. **Complexity of Multi-Agent Systems:** Modern applications involve multiple agents working together, such as in autonomous vehicles or distributed sensor networks. These systems require proper handling of uncertainty across all agents to make consistent and safe decisions.
5. **Advancements in Technology:** With the availability of modern machine learning frameworks, computational power, and simulation tools, it is now possible to design systems that can explicitly model and use uncertainty in decision-making.

The motivation behind Cognix is to develop a system that not only improves prediction accuracy but also makes the decision-making process more reliable, transparent, and suitable for real-world deployment.

1.3 Literature Survey

A substantial amount of research has been conducted in the areas of uncertainty estimation, probabilistic learning, and multi-agent systems. These studies provide a strong foundation but also reveal key limitations that motivate the development of Cognix.

Early work in probabilistic machine learning by Bishop [2] and Murphy [22] established the theoretical basis for modeling uncertainty in predictions. Bayesian Neural Networks, introduced by Neal [21], further extended these ideas by placing probability distributions over model parameters, enabling principled uncertainty estimation.

In deep learning, Gal and Ghahramani [3] proposed Monte Carlo Dropout as an approximation to Bayesian inference, while Lakshminarayanan et al. [4] introduced deep ensembles for improved predictive uncertainty. Although effective, these methods either approximate the posterior or require multiple models, limiting their efficiency in real-time systems.

Kendall and Gal [13] highlighted the importance of distinguishing between aleatoric and epistemic uncertainty, particularly in computer vision tasks. This distinction is critical for understanding whether uncertainty arises from the environment or the model itself. However, most implementations focus on single-agent systems and do not extend to collaborative environments.

Calibration techniques such as those discussed by Guo et al. [5] and conformal prediction methods introduced by Vovk et al. [7] provide statistical guarantees on prediction reliability. Recent advancements [26] further improve the applicability of conformal methods in modern machine learning systems.

Research in multi-agent systems [17, 19] explores coordination and decision-making across distributed agents. However, these approaches often assume deterministic predictions and do not incorporate uncertainty as a first-class component of the system.

From the existing literature, it is evident that while uncertainty estimation, calibration, and multi-agent collaboration have been studied individually, there is a lack of integrated systems that combine all three aspects.

Cognix addresses this gap by proposing a unified framework that performs real-time uncertainty decomposition, multi-agent belief aggregation, and calibrated decision-making, making it suitable for deployment in complex real-world environments.

Chapter 2

Problem Definition and Scope

2.1 Problem Statement

Artificial Intelligence systems are increasingly used in critical domains such as autonomous vehicles, healthcare diagnostics, financial decision-making, and multi-robot systems. While these systems have achieved high accuracy, a major limitation still exists — they often produce decisions without properly representing uncertainty. As a result, systems may take incorrect actions with high confidence, which can lead to serious real-world consequences.

Another key challenge arises in multi-agent environments, where multiple AI systems or sensors contribute to a shared decision. In such settings, there is no unified mechanism to quantify, combine, and interpret uncertainty across different agents. Existing approaches typically focus on single-model predictions and fail to address how uncertainty should be communicated, calibrated, and used for decision-making in a distributed system.

Therefore, there is a need to design a system that not only generates predictions but also provides reliable uncertainty estimates, combines information from multiple agents, ensures calibrated confidence, and supports explainable decision-making. The Cognix system aims to address this gap by developing a structured framework for uncertainty-aware multi-agent decision systems.

2.1.1 Goals and Objectives

- **Goals**

1. Develop Cognix as a unified multi-agent decision system that incorporates uncertainty awareness at every stage of processing.
2. Enable reliable and interpretable decision-making by integrating prediction, uncertainty estimation, calibration, and explanation within a single framework.

3. Improve safety and robustness of AI systems operating in real-world environments by reducing overconfident incorrect decisions.
4. Build a scalable architecture that supports multiple agents and heterogeneous data sources such as sensors, simulations, and real-world inputs.

• Objectives

1. **Framework Development:** Design and implement a complete 5-module Bayesian multi-agent system for end-to-end uncertainty quantification and decision-making.
2. **Uncertainty Modeling:** Decompose uncertainty into aleatoric (irreducible data noise) and epistemic (reducible model uncertainty) at the individual agent level.
3. **Confidence Calibration:** Provide mathematically guaranteed and reliable confidence intervals using conformal prediction techniques.
4. **Efficient Communication:** Enable intelligent and uncertainty-aware communication between agents using Graph Neural Networks (GNNs).
5. **Uncertainty Attribution:** Identify and quantify the contribution of each agent to overall system uncertainty using Shapley value-based methods.
6. **Multi-Domain Validation:** Evaluate system performance across real-world domains such as autonomous driving (CARLA), medical diagnosis (CheXpert), and drone swarm search rescue.
7. **Regulatory Compliance:** Ensure the system meets modern AI safety and transparency standards such as ISO 26262, EU AI Act 2.0, and FDA guidelines.

2.1.2 Assumptions and Scope

• Assumptions

- ◇ **Data Availability:** It is assumed that sufficient training and testing data is available through simulation environments and real-world datasets.
- ◇ **Sensor Inputs:** The system assumes that agents provide structured input data such as images, numerical readings, or extracted features suitable for machine learning models.
- ◇ **Computational Resources:** It is assumed that adequate computational resources are available for model training and real-time inference.
- ◇ **Model Generalization:** The system assumes that models trained on simulation data can generalize reasonably well to real-world scenarios.

• Scope

-
- ◇ The project focuses on developing a multi-module AI system that integrates uncertainty estimation, probabilistic reasoning, calibration, communication, and decision-making.
 - ◇ The system is designed for multi-agent environments such as autonomous vehicles, robotics, and distributed sensing systems.
 - ◇ Validation is carried out using simulation platforms and controlled experimental setups.
 - ◇ The project emphasizes system design, model implementation, and evaluation rather than large-scale industrial deployment.
 - ◇ The system outputs include predictions along with uncertainty measures and explanations for decision-making.
 - ◇ User interface development and large-scale production deployment are not included in the scope of this project.

By clearly defining the problem, goals, and scope, the Cognix system aims to provide a reliable and structured approach for uncertainty-aware decision-making in modern AI systems.

Chapter 3

Methodology

3.1 Methodology

Refer to Fig. 3.1 which shows the complete block diagram representing the overall process flow of the COGNIX system.

1. **Input Acquisition:**

The system receives input data or observations, which are distributed across multiple agents. Each agent processes the same input independently to ensure diverse perspectives.

2. **Independent Agent Processing:**

Each agent uses its own model to analyze the input data and generate predictions.

3. **Uncertainty Estimation:**

Along with predictions, each agent estimates uncertainty using probabilistic techniques, considering both aleatoric and epistemic uncertainty.

4. **Inter-Agent Communication:**

Agents share their predictions and uncertainty values to improve overall system understanding.

5. **Prediction Aggregation:**

The outputs from all agents are combined using a suitable aggregation method.

6. **Conformal Calibration Layer (Module 3):**

The aggregated prediction is passed through the conformal calibration layer to ensure reliable confidence estimation.

7. **Final Decision Output:**

The system produces the final prediction along with a confidence interval.




Figure 3.1: COGNIX - Overall System Block Diagram

3.1.1 Uncertainty Detection Engine (Module 1)

Refer to Fig. 3.2 which illustrates the architecture of the Uncertainty Detection Engine.

This module is responsible for estimating the reliability of predictions generated by each agent. It operates at the initial stage of the pipeline and produces uncertainty-aware outputs that are used in subsequent decision-making stages.

The system is based on a Bayesian Neural Network (BNN), which models uncertainty by treating network parameters as probability distributions instead of fixed values. This allows the system to generate both predictions and associated confidence measures.

Two types of uncertainty are considered:

- **Aleatoric Uncertainty:** Arises due to noise or randomness in input data.
- **Epistemic Uncertainty:** Arises due to limited knowledge of the model.

The detailed internal workflow, including prior modeling, variational learning, sampling, and uncertainty decomposition, is represented in Fig. 3.2.

The output of this module is a structured tuple of the form $(\mu, \sigma_{aleatoric}, \sigma_{epistemic})$, which is used for further aggregation and calibration.




Figure 3.2: COGNIX - Uncertainty Detection Engine

Chapter 4

Outcome

4.1 Type of Project

The COGNIX project can be categorized as both application-oriented and research-oriented, with a strong focus on improving prediction reliability through calibration techniques.

- **Application-Oriented:**

The system is designed to be used in real-world scenarios where reliable decision-making is important. The Conformal Calibration Layer ensures that predictions are accompanied by meaningful confidence intervals.

- **Research-Oriented:**

The project explores the application of conformal prediction in a multi-agent environment. Extending calibration from single models to multiple interacting agents adds a research dimension to the work.

- **AI and Machine Learning Domain:**

The project uses concepts from machine learning, uncertainty estimation, and statistical calibration to improve system performance.

- The main outcomes of the Conformal Calibration Layer are:

- **Improved Reliability:** Predictions are no longer overconfident and better reflect actual performance.
- **Confidence Guarantees:** The system maintains a predefined confidence level (e.g., 95%).
- **Better Decision-Making:** The system can identify uncertain cases and avoid risky decisions.

-
- **Enhanced Interpretability:** Users can understand how confident the system is about its predictions.

4.2 Expected Outcomes

4.2.1 Input and Output Format

The input to Module 1 is the raw sensor observation collected from an agent in the system. This may include data from different sources such as camera, lidar, or radar.

Input:

```
{
  agent_id:    3,
  observation: [sensor readings from camera / lidar / radar]
}
```

The Uncertainty Detection Engine processes this input using a Bayesian Neural Network to estimate both the prediction and its associated uncertainties. The system separates uncertainty into aleatoric (due to data noise) and epistemic (due to lack of knowledge).

The output of Module 1 is a structured uncertainty representation for each agent:

Output:

```
{
  agent_id:          3,
  prediction_mu:     0.72,
  sigma_aleatoric:   0.15,
  sigma_epistemic:   0.09,
  total_uncertainty: 0.24
}
```

This output is forwarded to the next modules in the Cognix system, where it is used for belief aggregation and decision-making. By providing separate uncertainty values, the module ensures that later stages can interpret predictions more effectively.

4.2.2 Expected Performance

The module is expected to achieve reliable uncertainty decomposition with high accuracy and low latency. The system aims to maintain decomposition accuracy above 85 percent while ensuring that inference time per agent remains within 5 milliseconds. It should also be capable of detecting out-of-distribution scenarios and maintaining stable predictions across multiple runs.

4.2.3 Key Outcomes

The main outcomes of Module 1 are as follows:

- **Structured Uncertainty Output:** The system provides clear separation of prediction, aleatoric uncertainty, and epistemic uncertainty for each agent.

-
- **Improved Decision Support:** By distinguishing between different types of uncertainty, the system enables better decision-making in uncertain situations.
 - **Real-Time Performance:** The module operates within strict time constraints, making it suitable for real-world applications.
 - **Better Interpretability:** The output allows users and downstream modules to understand the source of uncertainty in predictions.

Chapter 5

Project Plan

5.1 Project Timeline




Table 5.1: Project Timeline

5.2 Team Organisation

The Cognix project is developed by a team of four members, each responsible for specific modules of the complete system.

5.2.1 Team Structure

1. Komal — Module 1: Uncertainty Detection Engine

Komal is responsible for designing and implementing the Bayesian Neural Network-based uncertainty decomposition pipeline that forms the foundation of the entire Cognix system.

Key Responsibilities:

- ◊ Aleatoric and Epistemic Uncertainty Decomposition using BNN
- ◊ Variational Inference training with ELBO loss
- ◊ Posterior sampling pipeline (S=50 weight samples)
- ◊ Structured output tuple generation (μ, σ_a, σ_e)
- ◊ CARLA simulation scenario design for validation
- ◊ DuckieBot hardware latency testing and optimisation
- ◊ Literature review for uncertainty detection methods

2. Aryan — Module 2: Bayesian Belief Network

Aryan is responsible for multi-agent probabilistic belief fusion using Loopy Belief Propagation, weighted by epistemic uncertainty from Module 1.

Key Responsibilities:

- ◊ Bayesian Belief Network design and implementation
- ◊ Loopy Belief Propagation algorithm
- ◊ Epistemic-uncertainty-weighted message passing
- ◊ Convergence testing and collective posterior computation

3. Aditya — Module 3: Conformal Calibration Engine

Aditya is responsible for providing statistical guarantees on prediction reliability through conformal calibration techniques.

Key Responsibilities:

- ◊ Conformal Prediction implementation with coverage guarantees
- ◊ Calibration of uncertainty outputs from Module 1 and Module 2
- ◊ Two-stage hierarchical calibration (individual and collective)
- ◊ Non-conformity score computation and threshold selection

4. Zaid — Module 4 and 5: Communication and Decision Engine

Zaid is responsible for efficient inter-agent uncertainty communication and final risk-aware decision making with attribution.

Key Responsibilities:

- ◇ Graph Attention Network for uncertainty-aware communication
- ◇ Bandwidth-constrained message selection protocol
- ◇ Task-specific decision threshold design
- ◇ Shapley value-based uncertainty attribution per agent
- ◇ Escalation decision tree implementation
- ◇ Final decision output formatting and logging

Chapter 6

Experimental setup

In order to evaluate the performance of the COGNIX system, a controlled experimental setup was designed. The aim was to analyze not only prediction accuracy but also the reliability and calibration of outputs. The system was tested under different conditions to understand how well it handles uncertainty in a multi-agent environment. The following sections describe the dataset, technologies used, and performance parameters considered during evaluation.

6.1 Data Set

Synthetic and Benchmark Dataset

- **Source:** Combination of synthetically generated data and standard machine learning benchmark datasets.
- **Formats:** Structured numerical data along with simulated observations representing different agent inputs.
- **Content Variety:**
 - ◊ Regression-based data for continuous prediction tasks
 - ◊ Classification scenarios with probabilistic outputs
 - ◊ Simulated multi-agent observations with varying uncertainty levels

6.2 Technology Used

- **Machine Learning Framework:**
 - ◊ PyTorch for building and training predictive models
- **Uncertainty Estimation:**
 - ◊ Monte Carlo Dropout and ensemble-based techniques for estimating uncertainty
- **Calibration Layer:**

-
- ◊ Conformal Prediction implemented as a post-processing step for generating calibrated intervals
 - **Backend & Processing:**
 - ◊ Python 3.x for overall system implementation
 - ◊ NumPy and Pandas for data manipulation and processing
 - **Development Tools:**
 - ◊ Git for version control and collaboration
 - ◊ Jupyter Notebook for experimentation and analysis

6.3 Performance Parameters

To evaluate the effectiveness of the system, both prediction performance and calibration quality were measured. The following parameters were considered:

- **Prediction Accuracy:**
 - Measures how close the predicted values are to actual values
 - Used standard metrics such as Mean Squared Error (MSE) or accuracy for classification
- **Calibration Accuracy:**
 - Evaluates how well the predicted confidence matches actual correctness
 - Goal: Confidence levels should align with real-world outcomes
- **Coverage:**
 - Percentage of true values that lie within the predicted confidence interval
 - Target: Coverage close to predefined confidence level (e.g., 95%)
- **Prediction Interval Width:**
 - Measures the size of the confidence interval
 - Narrower intervals with correct coverage indicate better performance
- **Computation Time:**
 - Time required to generate predictions and calibration intervals
 - Evaluated to ensure system efficiency
- **System Reliability:**
 - Measures consistency of predictions across different agents
 - Helps in evaluating robustness of multi-agent decision-making

Chapter 7

PO & PSO Mapping

7.1 Program Outcome (PO) Mapping

PO	Description	Level	Justification
PO1	Engineering Knowledge: Apply knowledge of mathematics, science, and engineering fundamentals to AI and DS problems.	3	Cognix applies advanced mathematics including Bayesian inference, probability distributions, Variational Inference, ELBO loss, Shapley values, graph theory, and conformal prediction to solve complex multi-agent uncertainty quantification problems.
PO2	Problem Analysis: Identify, formulate, and analyze complex problems using principles of math, natural science, and engineering.	3	The core problem of over-confident AI systems was identified, formally defined, and analyzed through a review of three recent IEEE papers, reaching the substantiated conclusion that aleatoric and epistemic uncertainty must be separated in real-time multi-agent settings.

PO3	Design/Development of Solutions: Design system components or processes considering societal, safety, and environmental factors.	3	A complete five module system was designed addressing public safety, autonomous driving, medical diagnosis, financial trading with appropriate consideration for regulatory compliance including ISO 26262 and EU AI Act 2.0.
PO4	Conduct Investigations of Complex Problems: Conduct research based investigations, design experiments, analyze data, and synthesize conclusions.	3	Experiments are designed using CARLA simulator across four uncertainty scenario categories, with analysis of decomposition accuracy, ECE, latency, and OOD detection all grounded in research based investigation methodology.
PO5	Modern Tool Usage: Create and apply modern tools, modeling, and engineering techniques.	3	PyTorch 2.0.1, Pyro, BayesianTorch, CARLA 0.9.14, PettingZoo, DuckieBot, PyTorch Geometric, TorchScript, and Weights and Biases are all applied extensively across the Cognix system.
PO6	The Engineer and Society: Apply contextual knowledge to societal, health, safety, legal, and cultural issues.	3	Cognix directly addresses societal safety, autonomous vehicle accidents, medical AI misdiagnosis, financial flash crashes and aligns with legal regulatory requirements of ISO 26262, EU AI Act 2.0, and FDA 2026 guidelines.

PO7	Ethics: Apply ethical principles and commit to professional ethics and responsibilities.	3	Cognix is fundamentally motivated by the ethical responsibility of deploying AI in safety-critical domains. The system ensures AI is transparent and honest about its confidence, not overconfident before making decisions affecting human lives.
PO8	Individual and Team Work: Function effectively as an individual and in teams.	3	Cognix is developed by a four member team with each member owning one module functioning effectively both individually and collaboratively in a multi-disciplinary AI and robotics project.
PO9	Communication Skills: Communicate effectively on complex engineering activities.	3	Technical documentation, IEEE literature review, seminar report, and PPT presentation were all prepared, communicating complex uncertainty quantification concepts clearly to both technical and non-technical audiences.
PO10	Project Management and Finance: Apply project management and financial principles in engineering work.	3	The project was structured across eight phases with clear module-wise division of work, timelines, integration milestones, and hardware procurement planning for DuckieBot and CARLA infrastructure.

PO11	Life-long Learning: Recognize the need for life-long learning amidst technological change.	3	The team independently studied cutting edge research from 2022 to 2025 on Bayesian Neural Networks, conformal prediction, and graph-based uncertainty communication all topics beyond standard curriculum demonstrating self directed life long learning.
------	---	---	---

Table 7.1: Program Outcomes (PO) Mapping – Cognix

7.2 Program Specific Outcome (PSO) Mapping

PSO Number	Description	Mapping Level	Justification
PSO1	Interpret data and apply Artificial Intelligence algorithms to provide solutions.	3	Cognix interprets multi-modal sensor data from camera, LIDAR, and radar agents and applies Bayesian Neural Networks, Graph Neural Networks, Loopy Belief Propagation, and conformal prediction algorithms to provide safe and calibrated solutions for multi-agent decision making in autonomous driving, medical diagnosis, and distributed robotics.

PSO2	Apply standard practices, strategies and use appropriate models of data analytics to discover knowledge.	3	Standard Bayesian inference practices, Shapley value attribution analysis, and Expected Calibration Error evaluation strategies are applied to multi-agent sensor data to discover actionable knowledge — identifying which agents are reliable, which uncertainty sources can be reduced, and which scenarios require human intervention.
------	--	---	--

Table 7.2: Program Specific Outcomes (PSO) Mapping – Cognix

Chapter 8

Summary

Cognix is designed to solve a problem that existing AI systems largely ignore the inability to explain why they are uncertain and what should be done about it. The system operates across five interconnected modules, where each agent first decomposes its sensor observations into two distinct uncertainty components, fuses its beliefs with neighbouring agents through probabilistic message passing, receives a mathematically guaranteed calibration ensuring stated confidence levels are accurate, communicates only the most informative uncertainty signals to other agents through a learned graph-based protocol, and finally receives a complete decision output that not only recommends an action but also identifies which specific agent is the weakest link in the collective reasoning chain. Applied to autonomous driving, medical diagnosis, and multi-robot coordination, Cognix moves AI decision making away from blind confidence toward a model of structured and explainable self-awareness, where every prediction carries a justified level of trust rather than an arbitrary score.

References

- [1] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. MIT Press, 2016.
- [2] C. M. Bishop, *Pattern Recognition and Machine Learning*. Springer, 2006.
- [3] Y. Gal and Z. Ghahramani, "Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning," in ICML, 2016.
- [4] B. Lakshminarayanan, A. Pritzel, and C. Blundell, "Simple and Scalable Predictive Uncertainty Estimation using Deep Ensembles," in NeurIPS, 2017.
- [5] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On Calibration of Modern Neural Networks," in ICML, 2017.
- [6] J. Platt, "Probabilistic Outputs for Support Vector Machines," 1999.
- [7] V. Vovk, A. Gammerman, and G. Shafer, *Algorithmic Learning in a Random World*. Springer, 2005.
- [8] A. N. Angelopoulos and S. Bates, "A Gentle Introduction to Conformal Prediction and Distribution-Free Uncertainty Quantification," 2021.
- [9] Y. Romano, E. Patterson, and E. Candès, "Conformalized Quantile Regression," in NeurIPS, 2019.
- [10] J. Lei et al., "Distribution-Free Predictive Inference for Regression," JASA, 2018.
- [11] V. Kuleshov, N. Fenner, and S. Ermon, "Accurate Uncertainties for Deep Learning Using Calibrated Regression," in ICML, 2018.
- [12] M. Abdar et al., "A Review of Uncertainty Quantification in Deep Learning," Information Fusion, 2021.
- [13] A. Kendall and Y. Gal, "What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision?" in NeurIPS, 2017.
- [14] T. G. Dietterich, "Ensemble Methods in Machine Learning," in Multiple Classifier Systems, 2000.
- [15] Z.-H. Zhou, *Ensemble Methods: Foundations and Algorithms*. CRC Press, 2012.
- [16] L. Buşoniu, R. Babuška, and B. De Schutter, "A Comprehensive Survey of Multiagent Reinforcement Learning," IEEE TSMC, 2008.

-
- [17] F. A. Oliehoek and C. Amato, *A Concise Introduction to Decentralized POMDPs*. Springer, 2016.
 - [18] P. Stone and M. Veloso, "Multiagent Systems: A Survey from a Machine Learning Perspective," *Autonomous Robots*, 2000.
 - [19] K. Zhang et al., "Multi-Agent Reinforcement Learning: A Selective Overview," *IEEE TNNLS*, 2021.
 - [20] M. Sensoy, L. Kaplan, and M. Kandemir, "Evidential Deep Learning to Quantify Classification Uncertainty," in *NeurIPS*, 2018.
 - [21] R. M. Neal, *Bayesian Learning for Neural Networks*. Springer, 1995.
 - [22] K. P. Murphy, *Machine Learning: A Probabilistic Perspective*. MIT Press, 2012.
 - [23] C. M. Bishop, "Pattern Recognition and Machine Learning (Information Science and Statistics)," 2007.
 - [24] G. Shafer and V. Vovk, "A Tutorial on Conformal Prediction," *JMLR*, 2008.
 - [25] H. Papadopoulos et al., "Inductive Confidence Machines for Regression," *ECML*, 2002.
 - [26] A. Angelopoulos et al., "Conformal Prediction: A Unified Review," *Foundations and Trends in ML*, 2023.
 - [27] Y. Liu et al., "A Survey on Deep Learning-Based Uncertainty Estimation," 2020.