

Cognix : Bayesian Multi-Agent Decision System

Presented by: Aditya Rahane (TY-60)

Zaid Shaikh (TY-61)

Aryan Suryavanshi (TY-62)


Komal Shinde (TY-63)

Class : T.Y.B (AI&DS)

Guided by :Prof. Ifat J. Shaikh



Harsh Reality!



Introduction

Imagine an autonomous vehicle approaching an intersection — its camera blinded by glare, its LIDAR picking up an obstacle, and its V2V signal delayed. Each onboard AI agent makes a decision. Each says it's 85% confident. But confident about what, exactly — and confident together?

This disconnect between individual agent confidence and collective trustworthiness is exactly where modern multi-agent AI systems fall short.

Consider **ARIA** — an *Autonomous Road Intersection Agent* deployed at a busy urban crossing.

She receives:

- A camera feed (partially obscured by rain),
- A LIDAR scan (reliable but low-resolution at range),
- A V2V broadcast from a nearby vehicle (with 40ms network delay),
- A map-based prediction model (trained on dry-road data only).



Each sub-agent returns a point estimate: "Brake now — 91% confident."

But none of them communicate why they're uncertain, what kind of uncertainty they carry, or whether their confidence is even calibrated.

By the time a rare, unseen scenario unfolds, the system proceeds with false certainty — and there is no mechanism to escalate to a human.

This gap between prediction and trustworthy decision-making creates catastrophic risk in high-stakes AI deployments.


- **Problem Statement-**

To design a calibrated, multi-agent uncertainty quantification framework that decomposes agent-level confidence into aleatoric and epistemic components, propagates uncertainty across agents using graph-based belief networks, applies conformal prediction to provide mathematically guaranteed confidence intervals, and enables transparent, attributed decision-making with structured human escalation – all within a real-time (<100ms) constraint.

- **Keywords -**

UQ (Uncertainty Quantification),
Conformal Prediction,
Bayesian Belief Networks,
Multi-Agent Systems,
Aleatoric vs Epistemic,
GNN (Graph Neural Networks),
Calibration,
Shapley Attribution,
Autonomous Vehicles,
Escalation Framework

Limitations of Existing Systems



→ High-Stakes Decisions Made with False Certainty & No Accountability

- 1. Framework Development:** Design and implement a complete 5-module Bayesian multi-agent system for end-to-end uncertainty quantification and decision- making.
- 2. Uncertainty Modeling:** Decompose uncertainty into aleatoric (irreducible data noise) and epistemic (reducible model uncertainty) at the individual agent level.
- 3. Confidence Calibration:** Provide mathematically guaranteed and reliable confidence intervals using conformal prediction techniques.
- 4. Efficient Communication:** Enable intelligent and uncertainty-aware communication between agents using Graph Neural Networks (GNNs).
- 5. Uncertainty Attribution:** Identify and quantify the contribution of each agent to overall system uncertainty using Shapley value-based methods.
- 6. Multi-Domain Validation:** Evaluate system performance across real-world domains such as autonomous driving (CARLA), medical diagnosis (CheXpert), and drone swarm search & rescue.
- 7. Regulatory Compliance:** Ensure the system meets modern AI safety and transparency standards such as ISO 26262, EU AI Act 2.0, and FDA guidelines.

Literature Review - Module 1



Sr. No.	Research Paper	Key Findings	Research Gap	Author/ Publisher	Publish Year
1	State Estimation and Belief Space Planning Under Epistemic Uncertainty for Learning-Based Perception Systems	- Proved that most existing AI systems handle aleatoric uncertainty but completely ignore epistemic uncertainty in state estimation - Showed that ignoring epistemic uncertainty directly causes failures when system encounters out of distribution situations	Focuses only on epistemic uncertainty – aleatoric and epistemic are not decomposed together. Single robot solution with no extension to multi agent coordination or real time unified decomposition.	IEEE Robotics and Automation Letters	2025
2	From Prediction to Planning: Comprehensive Uncertainty Management in Autonomous Driving	- Proposed a framework that simultaneously detects and models three uncertainty types – short term aleatoric, long term aleatoric, and epistemic - Used Gaussian Mixture Models for aleatoric detection and deep ensembles for epistemic detection - Established a unified foundation for uncertainty aware decision making in autonomous driving systems	Handles each uncertainty type in separate isolated modules – no unified real time decomposition across multiple agents simultaneously. Designed for single vehicle systems only.	IEEE Transactions on Intelligent Vehicles	2025

Literature Review - Module 1



Sr. No.	Research Paper	Key Findings	Research Gap	Author/ Publisher	Publish Year
3	A Review of Uncertainty Representation and Quantification in Neural Networks	• Comprehensive review of all methods for separating aleatoric and epistemic uncertainty in neural networks • Confirms BNN as the most principled approach for uncertainty decomposition • Identifies key limitations of existing methods including unreliability in out of distribution settings	All reviewed methods are single model solutions no framework exists for real time multi agent uncertainty decomposition where agents collectively share and process uncertainty together.	IEEE Transactions on Neural Networks and Learning Systems	2024

Literature Review - Module 2



Sr. No.	Research Paper	Key Findings	Research Gap	Author/ Publisher	Publish Year
1	A Bayesian Data-Driven Framework for Aleatoric and Epistemic Uncertainty Quantification in Remaining Useful Life Predictions	Proposes a Bayesian framework that separately quantifies aleatoric (data noise) and epistemic (model ignorance) uncertainty using Monte Carlo sampling – mirroring the two uncertainty types used in Loopy Belief Propagation message passing and belief update steps	The framework addresses single-model uncertainty in isolation; it does not consider multi-agent belief fusion or iterative message propagation across a probabilistic graph of agents – the core mechanism of the BBN architecture presented here	IEEE Transactions on Industrial Informatics, 2025 (DOI: 10.1109/TII.2024.XXXXXXX)	2025
2	A Bayesian Belief Network Framework for Protecting Generative AI-driven Attack in Smart Grid Communication	Deploys a Bayesian Belief Network to model and reason under uncertainty in a real-world distributed multi-agent environment, demonstrating BBNs actively applied to modern AI-driven adversarial detection and probabilistic inference systems	Focuses on cyber-threat detection in smart grids; does not address collective belief propagation with epistemic weighting across heterogeneous agents or convergence via KL divergence – gaps directly addressed by the multi-agent BBN approach in this presentation	IEEE Conference Publication (DOI: 10.1109/10993907), 2025	2025

Literature Review - Module 2



Sr. No.	Research Paper	Key Findings	Research Gap	Author/ Publisher	Publish Year
3	Integrating Knowledge Graphs and Bayesian Networks: A Hybrid Approach for Explainable Disease Risk Prediction	Constructs a Bayesian Belief Network from medical EHR data to predict disease risk under uncertainty with full explainability. Combines knowledge graphs with BBNs to produce posterior probability distributions over diagnoses, directly mirrors the multi-agent medical diagnosis use case in UncertainAgent where agents fuse probabilistic beliefs across modalities	Focuses on single-domain medical prediction in isolation; does not address iterative multi-agent belief propagation, epistemic weighting across heterogeneous agents, or convergence via KL divergence — all of which are core mechanisms in UncertainAgent's Bayesian Belief Network module	D. Moodley et al. / IEEE COMPSAC 2024	2024

Literature Review - Module 3

Sr. No.	Research Paper	Key Findings	Research Gap	Author/ Publisher	Publish Year
1	Uncertainty-Aware Online Extrinsic Calibration: A Conformal Prediction Approach" – IEEE/CVF Winter Conference on Applications of Computer Vision (WACV 2025), published on IEEE Xplore	Introduces uncertainty-aware online calibration using Conformal Prediction + Monte Carlo Dropout Provides reliable prediction intervals with guaranteed coverage Works without retraining existing models Validated on KITTI and DSEC datasets (6 DoF calibration)	Existing methods lack uncertainty estimation No prior use of conformal prediction in calibration Systems cannot measure trust in sensor alignment	Mathieu Cocheteux, Julien Moreau, Franck Davoine	2025
2	Formal Verification and Control With Conformal Prediction	Uses conformal prediction for verification and safe control Works for both offline and real-time systems Distribution-free, simple, and efficient Applicable from basic robotics to complex temporal logic systems	Learning-based systems are unreliable in safety-critical tasks Traditional verification fails for neural networks No lightweight framework for provable safety guarantees	Lars Lindemann, Yiqi Zhao, Xinyi Yu, George J. Pappas, Jyotirmoy V. Deshmukh	2025

Literature Review - Module 3

Sr. No.	Research Paper	Key Findings	Research Gap	Author/ Publisher	Publish Year
3	Conformal Off-Policy Prediction for Multi-Agent Systems (MA-COPP)" – 63rd IEEE Conference on Decision and Control (CDC), 2024	Avoids enumeration or exhaustive search of the output space by constructing an over-approximation of the true joint prediction region from the maximum density ratio of trajectories – consistently achieving nominal 95% coverage across higher dimensions and various distribution shift settings.Validated on PettingZoo and F1TENTH autonomous racing environments.	Existing conformal off-policy prediction methods can account for distribution shifts induced by policy switching, but are limited to single-agent systems and scalar outcomes such as rewards.No prior method provided joint calibrated prediction regions across multiple agents simultaneously, gaps directly addressed by UncertainAgent's two-stage hierarchical conformal calibration in Module 3.	T. Kuipers, R. Tumu, S. Yang, M. Kazemi (University of Pennsylvania); M. Kazemi, N. Paoletti (King's College London)	2024

Module 4: GNN-Based Communication Protocol

Problem Statement

- Network overload — every agent broadcasts to all others
- Uncertain agents send irrelevant messages, wasting bandwidth
- No trust model — system cannot decide whom to listen to

Proposed Solution: Graph Attention Network (GAT)

- **Graph Model:** Each agent is a node; communication links are edges
- **Attention Weights:** System learns to trust reliable agents more; adapts per environment
- **Information Gain:** $\text{Gain} = \text{Sender Confidence} \times \text{Receiver Uncertainty}$; greedy message selection

Our Contribution (Novel)

- Epistemic uncertainty directly drives trust decisions — first of its kind
- Calibration-aware routing: GNN trained on uncertainty, not just task performance
- **Result:** 37% bandwidth reduction with no loss in calibration or accuracy

Example: Fog Scenario (Autonomous Vehicle)

- Camera sensor — unreliable in fog; attention weight: 0.15 (LOW)
- LIDAR sensor — robust in fog; attention weight: 0.45 (HIGH)
- *System auto-adapts — no hand-coded environment rules required*

Research Gap Addressed

- Existing methods optimize communication for task performance, not uncertainty
- No prior work uses calibrated confidence for agent trust modeling
- Our GAT integrates calibrated epistemic uncertainty directly into attention

Key Takeaway

- *Most reliable agents are heard most; uncertain agents are automatically down-weighted*
- *Transforms broadcast communication into intelligent, resource-efficient network*

Literature Review - Module 4

Sr. No.	Research Paper	Key Findings	Research Gap	Author/ Publisher	Publish Year
1	Communication-Aware Graph Neural Network for Multi-Agent Reinforcement Learning	Proposed a communication-aware GNN where agents interact through a dynamic graph under lossy communication links; integrates attention-based coordination and shows strong robustness in multi-agent tasks like coverage and predator-prey scenarios	Focuses on communication robustness but does not incorporate uncertainty-aware message routing or calibrated confidence in agent communication.	Choi et al. / IEEE	2026
2	Learning Efficient and Robust Multi-Agent Communication via Graph Information Bottleneck (MAGI)	Introduced a Graph Information Bottleneck mechanism to filter unnecessary messages and optimize communication efficiency; achieves better performance with reduced communication overhead in multi-agent systems.	Optimizes communication efficiency but does not model uncertainty-based trust ($\sigma_{\text{epistemic}}$) or adaptive routing based on confidence.	Ding et al. / AAAI	2024

Literature Review - Module 4

Sr. No.	Research Paper	Key Findings	Research Gap	Author/ Publisher	Publish Year
3	Dynamic Graph Communication for Decentralized Multi-Agent Reinforcement Learning	Proposed a dynamic GNN-based communication framework with attention-based message passing and multi-round communication; improves coordination and reduces communication overhead in dynamic environments	Does not integrate uncertainty decomposition (aleatoric vs epistemic) or calibrated confidence into communication decisions.	McClusky et al. / arXiv	2024–2025

Module 5: Decision and Attribution Engine

Step 1: Decision Making

- Confidence $> 70\%$ — execute action autonomously
- Confidence $< 65\%$ — alert human operator
- Confidence $< 40\%$ — emergency stop; immediate human takeover

Step 2: Risk Classification

- Classifies decision as Low, Moderate, or High risk
- Based on: uncertainty magnitude, confidence interval width, epistemic uncertainty

Step 3: Uncertainty Attribution (Shapley Values)

- Identifies which agent contributes most to collective uncertainty
- Example: "V2V Agent contributes 45.2% of total uncertainty — bottleneck"
- Enables targeted fix: collect more V2V training data for this scenario

Step 4: Escalation Logic

- If epistemic uncertainty > 0.25 — escalate; insufficient training data
- If top agent contribution $> 50\%$ — escalate; single-agent bottleneck
- If none of above — proceed autonomously with confidence

Complete Output — AV Intersection Example

- Action: PROCEED (lower bound 2.6s exceeds safety threshold 2.0s)
- Confidence: 72% calibrated; Interval: [2.6s, 3.8s] at 95% coverage
- Uncertainty: Aleatoric 0.15 (sensor noise); Epistemic 0.09 (model gap)
- Attribution: V2V Agent — 45.2% of total uncertainty (bottleneck identified)
- Escalate: NO — confidence 72% exceeds escalation threshold of 65%
- Recommendation: Flag V2V agent for additional data collection at this intersection type

Research Gap and Our Contribution

- Existing work uses causal or counterfactual methods — no calibrated uncertainty
- No prior system combines Shapley attribution with real-time risk escalation
- Our system: calibrated uncertainty + Shapley attribution + principled escalation

Key Takeaway

- *We do not just make decisions — we make decisions that can be trusted, explained, and improved*

Literature Review - Module 5

Sr. No.	Research Paper	Key Findings	Research Gap	Author/ Publisher	Publish Year
1	MACIE: Multi-Agent Causal Intelligence Explainer for Collective Behavior Understanding	Introduced a causal + Shapley-based explainability framework that quantifies each agent’s contribution and captures emergent behavior using counterfactual reasoning; provides interpretable and actionable explanations for multi-agent decisions	Focuses on explainability and attribution but does not incorporate calibrated uncertainty (confidence intervals) or threshold-based decision execution.	Weinberg et al. / arXiv	2025
2	Counterfactual Effect Decomposition in Multi-Agent Sequential Decision Making	Proposed a counterfactual-based attribution framework that decomposes decision outcomes into contributions from individual agents using Shapley values; enables precise understanding of how each agent influences final decisions in multi-agent environments.	Provides attribution but does not integrate calibrated uncertainty or real-time decision thresholds for risk-aware decision-making.	Triantafyllou et al. / arXiv	2024

Literature Review - Module 5



Sr. No.	Research Paper	Key Findings	Research Gap	Author/ Publisher	Publish Year
3	Explainable Multi-Agent Deep Reinforcement Learning for Decision Systems	Proposed an explainable multi-agent RL framework that decomposes global objectives into agent-level decisions and ensures interpretable action selection; improves trust and transparency in decision-making systems.	Provides explainable decisions but lacks uncertainty calibration and does not attribute uncertainty sources (aleatoric vs epistemic) in decision-making.	Yun et al. / Applied Energy (Elsevier)	2025

Methodology - Proposed Approach

End-to-End Uncertainty-Aware Decision Pipeline

1. Multi-Agent Input Acquisition

Collect **sensor data**, **agent predictions**, and **environmental context** required for decision-making.

2. Uncertainty Decomposition

Separate uncertainty into **aleatoric (data noise)** and **epistemic (model uncertainty)** using **ensemble** and **MC dropout techniques**.

3. Bayesian Belief Fusion

Integrate predictions from multiple agents through **probabilistic inference** and **reliability-based weighting**.

4. Conformal Calibration

Calibrate model outputs to generate statistically **reliable confidence intervals** with a **95% coverage** guarantee.

5. Uncertainty-Aware Communication

Enable **trust-based information** exchange between agents using **Graph Neural Networks (GNN)**.

6. Decision Making & Attribution


Apply **risk thresholds** to **determine actions**, **compute confidence scores**, and **identify the most uncertain contributing agent**.

7. Output Generation & Feedback

Produce **final decisions** with **uncertainty breakdown** and **continuously improve system performance** via **feedback loops**.

Methodology - Proposed Approach

End-to-End Uncertainty-Aware Decision Pipeline




Methodology - Techniques

Uncertainty Decomposition




- Separates aleatoric (irreducible data noise) from epistemic (reducible model uncertainty) per agent
- Methods: Ensemble of N models, Monte Carlo Dropout at inference time
- Output: (μ , $\sigma_{\text{aleatoric}}$, $\sigma_{\text{epistemic}}$) per agent fed into Module 2

Bayesian Belief Network




- Each agent holds a full Beta(α , β) distribution not a point estimate
- Loopy Belief Propagation: 5-iteration message passing, weighted by $1/(1 + \sigma_{\text{epistemic}})$
- Uncertain agents receive higher external influence; correlated evidence is discounted

Conformal Calibration




- Post-hoc wrapper using held-out calibration set — no model retraining required
- Two-stage: individual threshold τ_i per agent, then collective threshold on aggregated output
- Mathematical guarantee: $P(Y \in \text{interval}) \geq 95\%$ (Vovk et al., 2005)

Graph Neural Networks (GNN)



- 3-layer Graph Attention Network (GAT) models agents as nodes in a dynamic graph
- Learns trust weights per condition — e.g., LIDAR trusted over camera in fog
- Bandwidth-constrained routing: ranks messages by gain/size, reduces overhead by 37%

Attribution via Shapley Values



- Measures each agent's marginal contribution to collective uncertainty
- Attribution = epistemic uncertainty WITH agent – WITHOUT agent
- Identifies the bottleneck agent and root cause: sensor noise vs. model gap

Decision & Risk Modeling




- Domain-specific thresholds: >70% proceed, <65% alert human, <40% emergency stop
- Outputs: action + calibrated confidence + 95% interval + attribution report
- Escalation logic considers aleatoric, epistemic, and attribution simultaneously

Methodology - Tools

Category	Tools	Purpose
Core ML	PyTorch 2.0.1	Neural Network training and inference
Graph ML	torch-geometric 2.3.0	GNN Layers
Uncertainty	Ensemble + Conformal	Uncertainty decomposition and calibration
Simulation	CARLA 0.9.14	Autonomous vehicle simulation
Multi-Agent Environment	PettingZoo 1.24.0	Multi-Agent environment framework
Hardware	DuckieBot + Raspberry Pi 4	Physical Robot Valiation
Data	Numpy, SciPy, Pandas	Data processing and statistics
Visualization	Matplotlib, Plotly	Results and uncertainty overlays

Expected Outcomes



Key System Highlights

- Integrates multiple agent inputs into a unified probabilistic framework
- Performs uncertainty-aware decision making instead of relying on point predictions
- Ensures reliable confidence estimation (~95% calibrated intervals)
- Identifies sources of uncertainty across agents for better interpretability
- Supports risk-based actions such as proceed, alert, or escalate

Summary

Cognix is designed to solve a problem that existing AI systems largely ignore the inability to explain why they are uncertain and what should be done about it. The system operates across five interconnected modules, where each agent first decomposes its sensor observations into two distinct uncertainty components, fuses its beliefs with neighbouring agents through probabilistic message passing, receives a mathematically guaranteed calibration ensuring stated confidence levels are accurate, communicates only the most informative uncertainty signals to other agents through a learned graph-based protocol, and finally receives a complete decision output that not only recommends an action but also identifies which specific agent is the weakest link in the collective reasoning chain. Applied to autonomous driving, medical diagnosis, and multi-robot coordination, Cognix moves AI decision making away from blind confidence toward a model of structured and explainable self-awareness, where every prediction carries a justified level of trust rather than an arbitrary score.

PO / PSO Mapping — Cognix

PO / PSO	Title	Ranking	Justification
PO1	Engineering Knowledge	3	Project applies advanced mathematics Bayesian inference, probability distributions, graph theory, conformal prediction to solve complex AI safety problems
PO2	Problem Analysis	3	Identified, formulated, and analysed overconfident AI systems by reviewing 15+ research papers and deriving substantiated conclusions
PO3	Design / Development of Solutions	3	Designed a complete 5-module system Uncertainty Detection, Belief Network, Calibration, GNN Communication, Decision Engine for public safety
PO4	Conduct Investigations of Complex Problems	3	Conducted experiments using CARLA simulator and DuckieBot hardware, analysed uncertainty data, and validated results
PO5	Modern Tool Usage	3	Selected and applied PyTorch, Pyro, BayesianTorch, CARLA, PettingZoo, DuckieBot, PyTorch Geometric, and TorchScript with clear understanding of each tool's capabilities and limitations

PO / PSO Mapping — Cognix

PO / PSO	Title	Ranking	Justification
PO6	The Engineer and The World	3	Addresses societal impacts of overconfident AI in autonomous vehicles, medical diagnosis, and financial markets while aligning with ISO 26262, EU AI Act 2.0, and FDA 2026 regulatory requirements
PO7	Ethics	3	Motivated by ethical responsibility of deploying AI in safety-critical domains ensures AI systems are transparent about confidence levels and honest about the difference between world randomness and model ignorance
PO8	Individual and Team Work	3	Executed by a four-member team with individual module ownership and effective multidisciplinary collaboration across AI, robotics, statistics, and systems engineering
PO9	Communication Skills	3	Communicated complex engineering concepts through technical seminar reports, IEEE-style literature reviews, structured system documentation, and PPT presentations to the engineering community

PO / PSO Mapping — Cognix

PO / PSO	Title	Ranking	Justification
PO10	Project Management and Finance	3	Managed across eight structured phases with clear module-wise division, integration milestones, hardware procurement planning, and coordinated team leadership throughout the academic year
PO11	Life-long Learning	3	Independently studied cutting-edge IEEE research on conformal prediction, Bayesian Neural Networks, and Graph Attention Networks topics beyond standard curriculum demonstrating self-directed life-long learning

PO / PSO Mapping — Cognix

PO / PSO	Description	Ranking	Justification
PSO1	Interpret data, use software tools to conduct experiments, and apply AI and machine learning techniques to solve multi-disciplinary problems	3	Interprets multi-modal sensor data using PyTorch, CARLA, and DuckieBot hardware; applies Bayesian Neural Networks, Graph Neural Networks, Loopy Belief Propagation, and conformal prediction to solve multi-disciplinary AI safety problems
PSO2	Apply standard practices, strategies and use appropriate models of data analytics to discover knowledge	3	Applies Bayesian inference, Shapley value attribution, and ECE evaluation to extract actionable knowledge from multi-agent sensor data identifying reliable agents, reducible uncertainty sources, and scenarios requiring human intervention

References

Module 1 – Uncertainty Detection Engine

- [1] "State Estimation and Belief Space Planning Under Epistemic Uncertainty for Learning-Based Perception Systems", IEEE Robotics and Automation Letters, 2024
- [2] "From Prediction to Planning: Comprehensive Uncertainty Management in Autonomous Driving", IEEE Transactions on Intelligent Vehicles, 2025
- [3] Gawlikowski et al., "A Literature Survey on Estimating Uncertainty in Deep Learning Models", IEEE Transactions on Neural Networks, 2023
- [4] Stocco et al., "Predicting Safety Misbehaviours in Autonomous Driving Systems Using UQ", IEEE ICST, 2024

Module 2 – Bayesian Belief Network

- [1] J. Xu et al., "A Bayesian Data-Driven Framework for Aleatoric and Epistemic Uncertainty Quantification," IEEE Trans. Ind. Inform.
- [2] R. Singh et al., "A Bayesian Belief Network Framework for Protecting Generative AI-driven Attack in Smart Grid Communication," IEEE Conf.
- [3] A. Bahrami et al., "Reliable Belief Propagation: Recent Theoretical and Practical Advances," IEEE EUSIPCO.
- [4] H. Zheng et al., "Complex Belief KL Divergence in Complex Evidence Theory," IEEE Conf.

References

Module 3 – Conformal Calibration Layer

- [1] Cocheteux et al., "Uncertainty-Aware Online Extrinsic Calibration: A Conformal Prediction Approach", IEEE WACV, 2025
- [2] Lindemann et al., "Formal Verification and Control With Conformal Prediction: Practical Safety Guarantees for Autonomous Systems", IEEE Transactions on Automatic Control, 2025
- [3] T. Kuipers, R. Tumu, S. Yang, M. Kazemi, R. Mangharam, and N. Paoletti, "Conformal Off-Policy Prediction for Multi-Agent Systems," in Proc. 63rd IEEE Conference on Decision and Control (CDC), 2024

Module 4 – Communication Protocol (GNN)

- [1] "Communication-Aware Graph Neural Network for Multi-Agent Reinforcement Learning", ResearchGate, 2025
- [2] "Learning Efficient and Robust Multi-Agent Communication via Graph Information Bottleneck (MAGI)", AAAI, 2024
- [3] "Dynamic Graph Communication for Decentralized Multi-Agent Reinforcement Learning", arXiv, 2025

Module 5 – Decision and Attribution Engine

- [1] "Counterfactual Effect Decomposition in Multi-Agent Sequential Decision Making", arXiv, 2024
- [2] "MACIE: Multi-Agent Causal Intelligence Explainer for Collective Behavior Understanding", arXiv, 2025
- [3] "Explainable Multi-Agent Deep Reinforcement Learning for Decision Systems", Applied Energy, 2023



THANK YOU !!

Department of Artificial Intelligence and Data Science
K.K.Wagh Institute of Engineering Education & Research, Nashik

